

Load-Balancer

CS 249

Team #2

Rashmeet Khanuja, Anusha Vijay, Steven
Yen




A challenge in distributed system

- An application is deployed on multiple nodes that form a cluster.
- Appears as one to the user.
- The nodes in the cluster may or may not have the same computing capacity.



With different configurations and capabilities, how do we ensure minimized response times and reduce idle time of CPU?

Answer: Load Balancer



Role of a Load Balancer

- It is a software or hardware component that distributes the workload amongst the nodes.
- Ensures no single server is overloaded.
 - Evenly distributes the load on each server based on their capacity.
- Acts as a single point of entry for the end user
 - Complexities of redirection is abstracted from users.
- Increases reliability, efficiency, scalability, reusability, and availability.



Examples of Load Balancers

NetScaler

Resonate

NGINX Plus

Zen Load Balancer

Amazon ELB


BalanceNG

LVS

F-5 Big-IP Local Traffic Manager

Cisco Load Balancer

Barracuda Load Balancer ADC



Web Server without load balancer


No Load Balancing



- In this monolithic architecture, as the number of users trying to access the server simultaneously increases, chances of failing to service user requests with optimal response times increases.
- Users may experience slow load times or something worse.



To Avoid This..

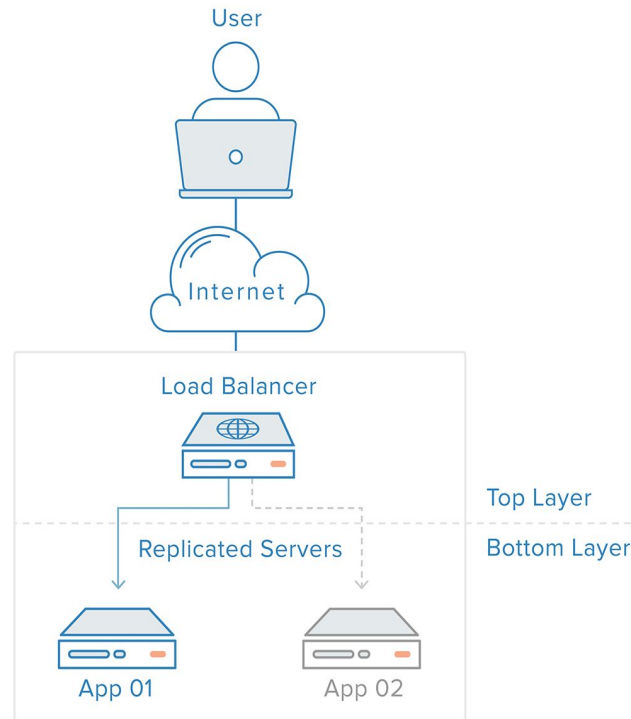


- A website with significant traffic will benefit from a distributed architecture with a load balancer

Distribute load between servers


To manage the incoming traffic:

- Load balancer listens on a port for incoming traffic from clients (e.g., a user's browser).
- Redirects any incoming client request (e.g., a HTTP GET request) to one of the many servers in the cluster.
- Server responds to the load balancer which is redirected back to the client.
- This is typically a software load balancer
- Benefits: hides the topology of the network/cluster from the client, providing some security. Faster response time for each client.



Software Load-Balancer

- Installed on all serving nodes
- Implement a combination of scheduling algorithms such as **Round-Robin** and **Least Connection First**
- **Weighted round-robin and weighted least connection.**
 - Useful when the servers/nodes have different capabilities (e.g. processing power/RAM).
 - Weight associated with the server depends on the hardware capabilities





Drawbacks

- Could be sensitive to OS versions
- Virtual appliance deployments could experience hypervisor dependencies
- Server's RAM and CPU





Hardware Load-Balancer

- Specialized routers or switches deployed between servers and clients
- Rely on firmware to supply the internal code base - the program
- Routing is either randomized (e.g., round-robin), or based on factors such as
 - available server connections,
 - server processing power,
 - resource utilization
- Also known as *layer 4-7 router* since they are implemented on Transport layer and Application layer of OSI model



Drawbacks

- Cost becomes considerable when global server load balancing or cross data center load balancing needs to be done
 - At least one appliance in each of the data centers
- A central box to manage load distribution between those appliances

